Distribution Plan

Currently, our database runs on MySQL hosted by Amazon RDS. We chose these services because of our familiarity with Amazon's offerings and MySQL. However, as our user base and demands grow, we will need to scale our infrastructure to ensure our application can handle the increased load efficiently.

Our plan is to transition to a more distributed server architecture. Instead of relying on a single server, we will distribute our primary server across multiple Amazon Web Services (AWS) EC2 instances worldwide. By leveraging AWS's global infrastructure, we can boost our computing power and ensure our service is fast and accessible globally.

Distributing our servers will make our system more flexible, resilient, and capable of handling a larger user base while maintaining smooth operations, even during peak traffic periods. AWS's robust infrastructure gives us confidence in making this transition seamlessly.

To handle high traffic and ensure high availability, we can configure auto-scaling and load balancing using AWS services like Elastic Load Balancing (ELB) and Auto Scaling Groups. Additionally, we can implement backup servers in different regions to mitigate the risk of physical server destruction, further enhancing our system's resilience. This means that one server failing will not make our application fail entirely as we are relying on a distributed network of servers not just one.

By implementing a distributed server architecture, we position ourselves to effectively scale our operations, provide a seamless user experience, and ensure reliable service delivery, even in the face of rapidly growing user demands or potential disruptions.